

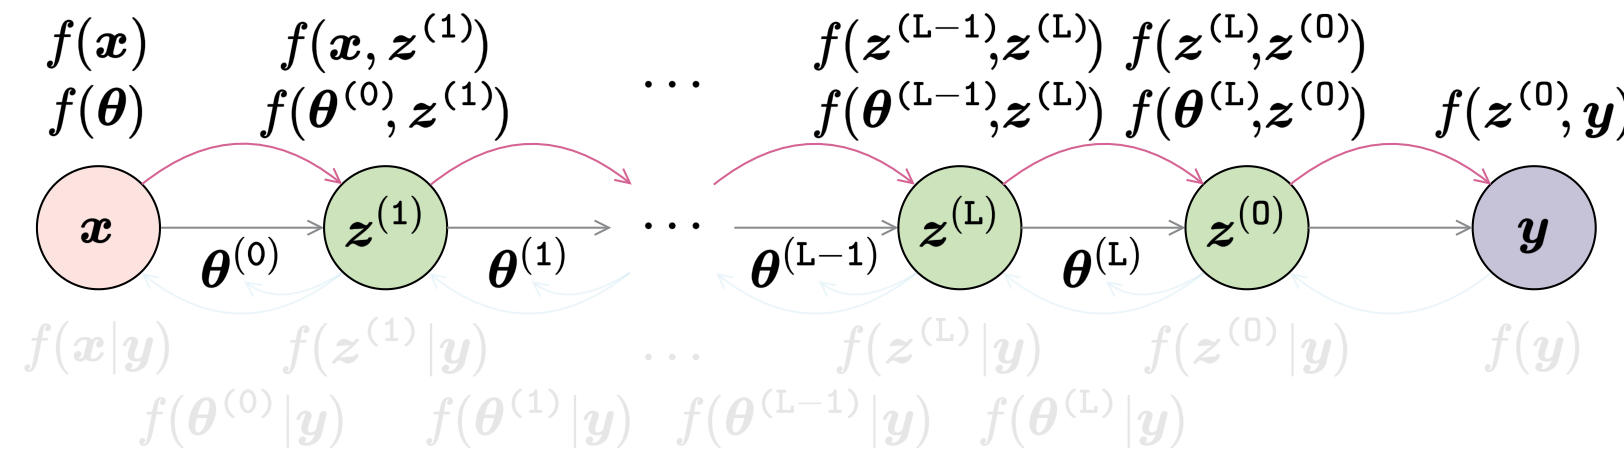
Implementation of a probabilistic Softmax for BNNs (Group 11)

David Wardan, Khalil Sabri, and Miquel Florensa – Polytechnique Montréal, Canada

(Adapted from the work proposed by Nguyen & Goulet)

Tractable Approximate Gaussian Inference (TAGI)

TAGI - Forward uncertainty propagation

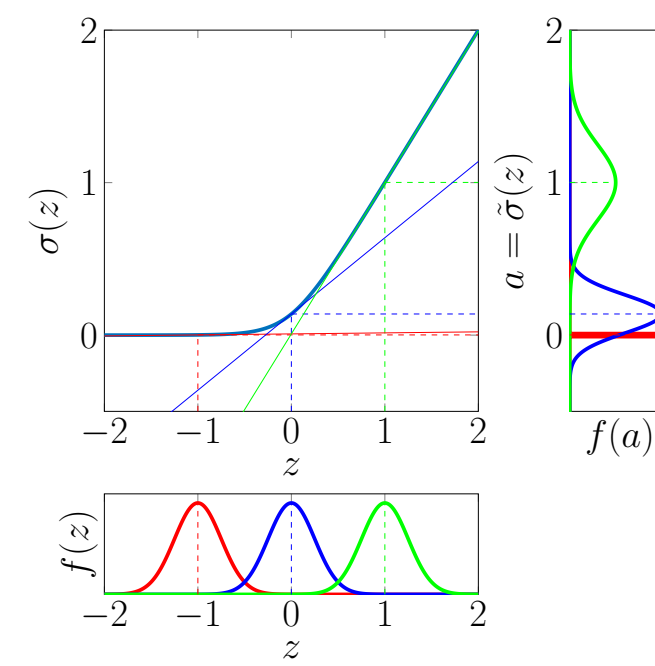
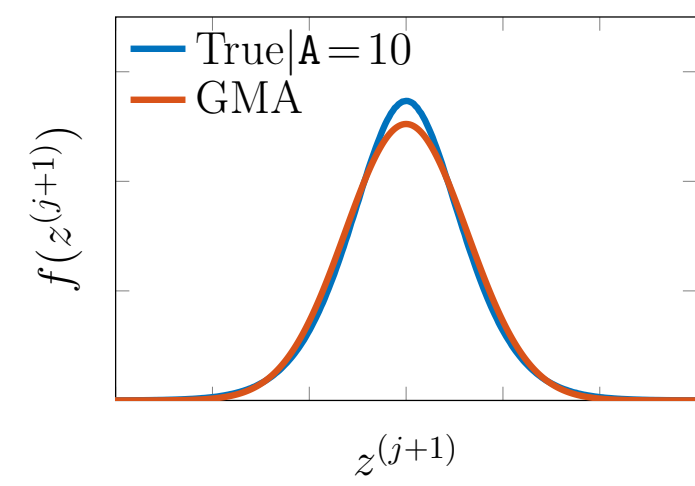


Challenge 1: GMA

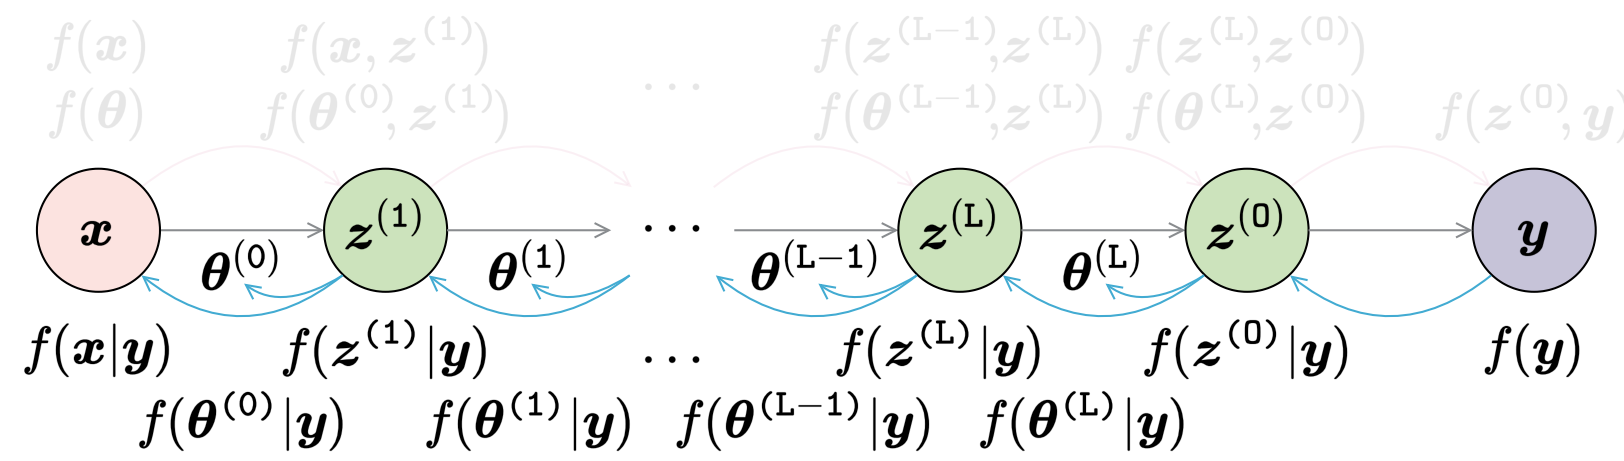
Challenge 2: Locally Linearize

$$\begin{aligned} z^{(1)} &= w^{(0)}x + b^{(0)}, \\ a^{(i)} &= \tilde{\sigma}(z^{(i)}), \\ z^{(i)} &= w^{(i)}a^{(i-1)} + b^{(i)}, \\ z^{(0)} &= w^{(L)}a^{(L)} + w^{(L)} \\ y &= z^{(0)} + v. \end{aligned}$$

where v are the observation noise.



TAGI - Backward parameter and hidden state inference



First, given the observations y , the output layer is updated as


$$\begin{aligned} f(z^{(0)} | y) &= \mathcal{N}(z^{(0)}; \mu_{Z^{(0)}|y}, \Sigma_{Z^{(0)}|y}) \\ \mu_{Z^{(0)}|y} &= \mu_{Z^{(0)}} + \Sigma_{YZ^{(0)}}^T \Sigma_Y^{-1} (y - \mu_Y) \\ \Sigma_{Z^{(0)}|y} &= \Sigma_{Z^{(0)}} - \Sigma_{YZ^{(0)}}^T \Sigma_Y^{-1} \Sigma_{YZ^{(0)}}. \end{aligned}$$

Next, the hidden units and parameters of the layer i^{th} are updated using a layer-wise procedure as

$$\begin{aligned} f(z | y) &= \mathcal{N}(z; \mu_{Z|y}, \Sigma_{Z|y}) \\ \mu_{Z|y} &= \mu_Z + J_Z (\mu_{Z^+|y} - \mu_{Z^+}) \\ \Sigma_{Z|y} &= \Sigma_Z + J_Z (\Sigma_{Z^+|y} - \Sigma_{Z^+}) J_Z^T \\ J_Z &= \Sigma_{ZZ^+} \Sigma_{Z^+}^{-1}, \end{aligned} \quad \begin{aligned} f(\theta | y) &= \mathcal{N}(\theta; \mu_{\theta|y}, \Sigma_{\theta|y}) \\ \mu_{\theta|y} &= \mu_{\theta} + J_{\theta} (\mu_{Z^+|y} - \mu_{Z^+}) \\ \Sigma_{\theta|y} &= \Sigma_{\theta} + J_{\theta} (\Sigma_{Z^+|y} - \Sigma_{Z^+}) J_{\theta}^T \\ J_{\theta} &= \Sigma_{\theta Z^+} \Sigma_{Z^+}^{-1}, \end{aligned}$$

where $\{\theta^+, Z^+\}$ is the short-hand notation of $\{\theta^{(j+1)}, Z^{(j+1)}\}$, and $\{\theta, Z\}$ is the short-hand notation of $\{\theta^{(j)}, Z^{(j)}\}$.

Regression toy example



- We re-implement the toy example presented by the authors using the numpy library in Python.
- This includes defining a fully connected neural network with one layer and 40 hidden units.
- The figure shows the output of the model on the test set with the epistemic uncertainty.

COLUMN 2 (Remax)



Key Results: TAGI-HRC vs. TAGI-Remax vs. Torch

Framework	Avg. Training Error	Avg. Test Error
Torch	10.05	10.69
TAGI-HRC	7.07	7.73
TAGI-Remax	7.14	7.78

Important Findings

- **Batch Size:** TAGI excels with smaller batches, while Torch performs better with larger batches.
- **Network Depth:** TAGI maintains stability as depth increases, unlike Torch, which struggles without batch normalization.
- **Effect of σ_v :** Higher σ_v values enhance TAGI's performance, especially in deeper networks.
- **Capacity Effect:** TAGI benefits more from increasing neurons/channels, showing improved generalization, while Torch fails to utilize additional capacity effectively.
- **Comparison of Implementations:** TAGI-HRC achieves better error rates overall, while TAGI-Remax shows higher stability but less frequent best-case performance.
- **Batch Normalization:** TAGI performs well even without batch normalization, handling deeper architectures (5+ layers) where Torch often diverges.

Ablation Study: Performance Comparisons and Capacity Effects



Test Error vs Network Depth: TAGI maintains low error rates, while Torch degrades significantly.

Impact of σ_v : Higher σ_v values improve TAGI's test error.

	Small Batch (16)		Large Batch (512)	
Framework	With BN	W.o BN	With BN	W.o BN
Torch	1.35	31.20	1.06	16.22
TAGI-HRC	1.97	1.74	4.11	15.99
TAGI-Remax	12.05	7.30	5.75	5.43

	Small (32)	Large (512)
Framework	Neurons/layer	Neurons/layer
Torch	7.03	7.32
TAGI-HRC	13.58	8.98
TAGI-Remax	9.25	6.91

py/cuTAGI – Open-source Bayesian deep-learning framework



github.com/lhnguyen102/cuTAGI

pip install pytagi

- **Performance-Oriented Kernels** written in C++/CUDA from scratch, with pybind11 for a seamless integration. It allows running on CPU & CUDA devices through a Python API.
- **Broad Architecture Support** of the basic DNN layer including *Linear*, *CNNs*, *Transposed CNNs*, *LSTM*, *Average pooling*, *Batch* and *Layer normalization*, enabling the building of mainstream architectures such as *Autoencoders*, *Transformers*, *Diffusion Models*, and *GANs*.
- **Model Building and Execution** currently supports sequential model building, with plans to introduce Eager Execution
- **Open Platform** providing access to its entire codebase.

TAGI-related references



- *Coupling LSTM Neural Networks and SSM through Analytically Tractable Inference*, (Vuong, Nguyen & Goulet, International Journal of Forecasting, 2024)
- *Analytically tractable hidden-states inference in Bayesian neural networks* (Nguyen and Goulet. Journal-to-conference track, ICLR 2024)
- *Analytically tractable heteroscedastic uncertainty quantification in Bayesian neural networks for regression tasks* (Deka, Nguyen & Goulet. Neurocomputing, 2024)
- *Tractable approximate Gaussian inference for Bayesian neural networks* (Goulet, Nguyen, & Amiri, JMLR, 2021)